

kg

Knowledge Graph

Reference Documentation

JSONL

Source of Truth

SQLite

Derived Index

MCP

LLM Integration

Lightweight · Git-friendly · Hybrid Search · Local-first

Contents

1	Introduction	1
1.1	Why kg?	1
1.2	Key Concepts	1
2	Architecture	2
2.1	Component Overview	2
2.2	Write Path	2
2.3	Read Path	2
2.4	Storage Layout	3
3	Installation	4
3.1	From Git (Latest)	4
3.2	Development / Editable	4
3.3	Optional Extras	4
4	Quickstart	5
4.1	Initialize a Project	5
4.2	First Notes	5
4.3	Search	5
5	CLI Reference	6
5.1	Core Commands	6
5.2	Node Commands	6
5.3	Search & Retrieval	6
5.4	Review & Quality	7
6	Search & Context	8
6.1	The Hybrid Search Pipeline	8
6.2	FTS Search	8
6.3	Vector Search	8
6.4	Score Calibration	9
6.5	Cross-Encoder Reranking	9
7	Configuration	10
7.1	Source File Indexing	10
8	Review & Token Budget	12
8.1	The Review System	12
8.2	Vote Quality Signals	12
9	MCP Server & Claude Code	13
9.1	Overview	13
9.2	MCP Tools	13
9.3	.claude → .kg Symlink	13
9.4	Session Context Hook	14
9.5	Stop Hook: Fleeting Notes Extraction	14
10	Web Viewer	15
10.1	Overview	15
10.2	Routes	15
10.3	Features	15

11 Troubleshooting	16
11.1 SQLite Disk I/O Error	16
11.2 Zero Files Indexed from Sources	16
11.3 Embeddings Coverage Below 100%	16
11.4 Vec Scores Missing from Calibration	16
12 Glossary	17

Introduction

kg is a lightweight, file-first knowledge graph designed for software projects and AI-assisted workflows. It keeps knowledge as human-readable JSONL files that are trivially grep-able, diff-able, and git-trackable, while deriving a SQLite index for fast hybrid search.

i Design Philosophy

Files are the source of truth. SQLite is a derived cache you can always rebuild with `kg reindex`. The graph survives any database corruption — your notes never will.

Why kg?

Traditional note systems force you to choose between structure (databases, heavy editors) and simplicity (plain text). kg offers a third path:

Plain text wins for durability

- `jq`, `rg`, `git log` work natively
- Merge conflicts are readable
- No lock-in: export is just `cat`
- Works offline, always

Index wins for retrieval

- FTS5 BM25 keyword search
- Vector similarity (local or cloud)
- Calibrated hybrid score fusion
- Cross-encoder reranking

Key Concepts

A **node** groups related **bullets** under a **slug**. Think of nodes as Zettelkasten cards and bullets as atomic facts on those cards.

```
project/
  .kg/
    nodes/
      my-topic/
        node.jsonl   ← bullets (append-only log, git-tracked)
        meta.jsonl   ← votes, usage counts (git-tracked)
      index/
        graph.db     ← SQLite: FTS5 + embeddings + backlinks (gitignored)
      kg.toml       ← project config
```

⚠ graph.db is gitignored

Never commit `graph.db`. It is always derivable from the JSONL files via `kg reindex`. The `.kg/nodes/` directory is what you version-control.

Architecture

Component Overview

The diagram below shows the main components and data flows in kg.




Figure 1: kg component architecture and data flows

Write Path

When you run `kg add my-topic "some fact"`:

1. The CLI appends a new bullet (JSONL line with stable UUID) to `.kg/nodes/my-topic/node.jsonl`, creating the directory if needed.
2. The inotify watcher detects the file change within milliseconds.
3. The watcher calls the indexer, which updates the FTS5 table in SQLite.
4. If embeddings are configured, the embedder generates a vector and stores it in the `embeddings` table.
5. Backlinks (`[[other-slug]]` references) are extracted and written to the `backlinks` table.


🔥 Only the watcher writes to SQLite

CLI commands like `kg add`, `kg update`, `kg delete` only write JSONL. The watcher is the **sole** writer to `graph.db`. This eliminates concurrent-write corruption.

Read Path

```
kg context "query"
```





Storage Layout

Path	Format	Purpose
kg.toml	TOML	Project configuration (name, embeddings, sources)
.kg/nodes/<slug>/node.jsonl	JSONL	Bullets — source of truth, git-tracked
.kg/nodes/<slug>/meta.jsonl	JSONL	Vote counts, usage stats (git-tracked)
.kg/index/graph.db	SQLite	FTS5 index, embeddings, backlinks (gitignored)
.kg/skills/	Markdown	Claude Code skills, via .claude→.kg symlink
~/.cache/kg/embeddings/	diskcache	Embedding vector cache (cross-project)

Table 1: File storage layout

Installation

From Git (Latest)

```
uv tool install "git+https://github.com/sasha-s/kg.git"
```

bash

With optional extras for embedding support and live file watching:

```
uv tool install "kg[embeddings,watch] @ git+https://github.com/sasha-s/kg.git"
```

bash

Development / Editable

```
git clone https://github.com/sasha-s/kg.git
```

bash

```
uv tool install --editable ./kg
```

Optional Extras

Extra	Packages	Use
embeddings	fastembed, google-genai, diskcache, numpy	Local + cloud embeddings
watch	inotify-simple	Live file watching (Linux)
turso	libsql	Turso remote SQLite (requires cmake)
dev	ruff, basedpyright, pytest	Development tools

Table 2: Optional install extras

Quickstart

Initialize a Project

```
# In your project directory:  
kg init          # writes kg.toml, creates .kg/  
kg start         # index + calibrate + watcher + vector server + MCP + hooks
```

After `kg start`, the `.claude → .kg` symlink is created so Claude Code can discover the MCP server and local skills automatically.

First Notes

```
kg add my-topic "discovered: X is faster than Y in this benchmark"  
kg add my-topic "gotcha: Y breaks when input is empty" --type gotcha  
kg add my-topic "decision: use X for production" --type decision
```

Search

```
kg search "fast benchmark"          # FTS keyword search  
kg context "which is faster"       # hybrid FTS + vector, calibrated + reranked
```

CLI Reference

Core Commands

Command	Description
kg init	Create kg.toml and .kg/ directory structure
kg start	Index + calibrate + start watcher + vector server + MCP + install hooks
kg stop	Stop watcher and vector server
kg reload	Hot-reload kg.toml config (sends SIGHUP to watcher, no restart)
kg reindex	Rebuild SQLite from all node.jsonl files (stops/restarts watcher)
kg calibrate	Compute FTS/vector score quantiles for hybrid fusion
kg upgrade	Run schema migrations + reindex (safe to run anytime)
kg status	Show node counts, calibration state, watcher, vector server, config
kg serve	Start MCP server on stdio (used by Claude Code)
kg web	Start web viewer at http://localhost:7343

Table 3: Core management commands

Node Commands

Command	Description
kg add <slug> <text> [--type TYPE]	Add a bullet to a node (auto-creates node if missing)
kg create <slug> <title> [--type TYPE]	Create a node with explicit title (idempotent)
kg show <slug>	Print all bullets in a node
kg update <bullet-id> <text>	Update bullet text by ID
kg delete <bullet-id>	Delete a bullet by ID
kg nodes [PATTERN]	List nodes (glob on slug, e.g. notes-*)
kg nodes --bullets	Sort nodes by bullet count descending
kg nodes --recent	Sort nodes by most-recently-updated
kg nodes --docs	Show _doc-* source-file nodes instead of curated nodes

Table 4: Node management commands

The --type flag accepts: fact, gotcha, decision, task, note, success, failure.

Search & Retrieval

Command	Description
<code>kg search <query> [-n N]</code>	FTS keyword search, returns raw results
<code>kg context <query> [-s SESSION_ID]</code>	Hybrid search + rerank, formatted for LLM injection
<code>kg context <query> -q <intent></code>	Use a separate query string for the reranking step

Table 5: Search and retrieval commands

Review & Quality

Command	Description
<code>kg review</code>	List nodes ordered by credits-per-bullet (review debt)
<code>kg review <slug></code>	Mark node as reviewed (clears token budget to 0)
<code>kg vote useful <bullet-id>...</code>	Signal that a bullet is high quality
<code>kg vote harmful <bullet-id>...</code>	Signal that a bullet is wrong or misleading

Table 6: Review and quality commands

Search & Context

The Hybrid Search Pipeline




Figure 2: Hybrid search and context retrieval pipeline

FTS Search

SQLite FTS5 BM25 ranking with automatic query expansion:

- Each search term becomes `term OR term*` (prefix wildcard for partial matches)
- Multiple terms are OR-joined for broader recall
- Weight controlled by `fts_weight` in `kg.toml` (default 0.5)

Vector Search

A lightweight HTTP vector server (port 7344) loads embeddings from SQLite at startup and serves approximate nearest-neighbor queries.

Prefix in kg.toml	Provider	Notes
gemini:...	Google Gemini	Cloud — requires GEMINI_API_KEY
openai:...	OpenAI	Cloud — requires OPENAI_API_KEY
fastembed:...	fastembed	Local, no API key — runs on-device
(bare model name)	fastembed	e.g. BAAI/bge-small-en-v1.5

Table 7: Supported embedding providers

Score Calibration

Raw FTS and vector scores are incomparable across queries and models. [calibration](#) maps them to percentile quantiles before fusion:

1. Sample 200 random nodes from the graph
2. Run FTS and vector search against each sample
3. Compute breakpoints at p0, p10, p25, p50, p75, p90, p100
4. At query time, binary-search each raw score into its quantile bucket
5. Fuse: $\text{score} = \text{fts_weight} \times \text{fts_q} + \text{vec_weight} \times \text{vec_q} + \text{dual_match_bonus}$

🔥 Tip

Run `kg calibrate` after adding many new nodes or after changing your embedding model. The watcher auto-recalibrates when more than `auto_calibrate_threshold` (5% by default) of bullets change.

Cross-Encoder Reranking

After fusion, the top-N candidates are passed to a cross-encoder [reranker](#). The reranker sees the full (query, bullet) pair — capturing semantic nuance that cosine similarity misses.

```
[search]
use_reranker = true
reranker_model = "Xenova/ms-marco-MiniLM-L-6-v2"
```

toml

Configuration

`kg.toml` lives in your project root. All sections are optional.

```
[kg] toml
name = "my-project"

[embeddings]
model = "gemini:gemini-embedding-001"
# model = "fastemb:BAAI/bge-small-en-v1.5"    # local alternative

[search]
fts_weight = 0.5
vector_weight = 0.5
dual_match_bonus = 0.1
use_reranker = true
reranker_model = "Xenova/ms-marco-MiniLM-L-6-v2"
auto_calibrate_threshold = 0.05  # recalibrate when 5% of bullets change

[review]
budget_threshold = 3000    # credits-per-bullet before review flag (~15 serves)

[server]
port = 7343      # web viewer
vector_port = 7344 # vector server

[[sources]]
name = "workspace"
path = "."
include = ["**/*.py", "**/*.md"]
use_git = true      # only index git-tracked files
```

API secrets in `.env` (gitignored by `kg init`):

```
GEMINI_API_KEY=...
TURSO_URL=libsql://...
TURSO_TOKEN=... bash
```

Source File Indexing

The `[[sources]]` section indexes source files into the same FTS index as curated nodes. Content is automatically chunked (avg 1500 bytes) without LLM extraction.

```
[[sources]] toml
name = "codebase"
```

```
path = "/path/to/project"
include = ["**/*.py", "**/*.ts", "**/*.md"]
exclude = [ ".kg/**", "node_modules/**"]
use_git = true
```

⚠ fnmatch differences: Python 3.11 vs 3.12

In Python 3.11, `fnmatch` treats `*` as matching `/` — so `*.py` matches `src/main.py`. In Python 3.12+, `*` does **not** match `/`. Use `**/*.py` patterns for reliability across Python versions.

Indexed files appear as `_doc-*` nodes in SQLite, browsable via `kg nodes --docs` and the web viewer.

Review & Token Budget

The Review System

Every time a node contributes bullets to kg context output, the character count is added to that node's **token budget**. When `budget ÷ bullet_count` exceeds `budget_threshold` (default 3000), the node is flagged for review.

Dividing by bullet count means every node — large or small — is held to the same standard.

```
kg review           # list flagged nodes, ordered by credits-per-  
bullet                                bash  
kg review <slug>        # inspect flagged bullets + mark as reviewed (clears  
budget)
```




Figure 3: Token budget lifecycle and review loop

Vote Quality Signals

Bullets can be voted on to signal quality. A bullet with more harmful than useful votes is shown with a Δ prefix in kg show output.

```
kg vote useful <bullet-id>    # signal high quality  
kg vote harmful <bullet-id>    # signal wrong or misleading
```

`kg review` surfaces nodes with net-harmful bullets even if under the budget threshold.

MCP Server & Claude Code

Overview

kg implements the Model Context Protocol (MCP) so any MCP-compatible LLM client can use it as a knowledge tool.

```
kg start # registers MCP automatically in ~/.claude/settings.json
```

bash

Or add manually to `~/.claude/settings.json`:

```
{
  "mcpServers": {
    "kg": {
      "command": "kg",
      "args": ["serve"],
      "cwd": "/path/to/your/project"
    }
  }
}
```

json

MCP Tools

Tool	Description
<code>memory_context(query, session_id?)</code>	Hybrid search + rerank, auto-dedup within session, LLM-ready output
<code>memory_search(query, limit?)</code>	Raw FTS keyword search, returns matching bullets
<code>memory_show(slug)</code>	Show all bullets for a node (read-only — does not clear budget)
<code>memory_add_bullet(node_slug, text, bullet_type?)</code>	Add a bullet to a node (auto-creates node if missing)
<code>memory_mark_reviewed(slug)</code>	Clear a node's token budget (marks node as reviewed)

Table 8: MCP tools exposed by `kg serve`

.claude → .kg Symlink

`kg start` creates `<project>/claude` → `.kg`. This lets Claude Code automatically discover:

- The MCP server registration (`settings.json`)
- Local skills under `.kg/skills/` (accessible as `.claude/skills/`)
- Project-specific instructions in `.kg/CLAUDE.md` (if present)

🔥 Bundled Skills

The `/note` and `/notes` skills ship with kg and are installed by `kg start`. `/note` adds a timestamped bullet with async cross-reference enrichment. `/notes` lists recent daily notes.

Session Context Hook

`kg start` installs a `UserPromptSubmit` hook that injects the current session ID into each prompt:

```
{
  "hooks": {
    "UserPromptSubmit": [
      {
        "hooks": [
          {
            "type": "command",
            "command": "python -m
kg.hooks.session_context"
          }
        ]
      }
    ]
  }
}
```

json

This enables per-session deduplication: bullets already surfaced in the current conversation are filtered from future `memory_context` results.

Stop Hook: Fleeting Notes Extraction

The stop hook runs at session end and promotes fleeting notes to permanent nodes:

1. Reads `_fleeting-<session-id>` node
2. Identifies patterns worth promoting to permanent nodes
3. Adds bullets to appropriate nodes with `[[cross-references]]`
4. Deletes the fleeting bullets after promotion

```
# During a session – low-friction scratch notes:
kg add _fleeting-abc123 "discovered: lock file blocks concurrent kg reindex"
```

bash

```
# At session end, the stop hook promotes this to the right permanent node.
```

Web Viewer

Overview

kg web serves a read-only UI at `http://localhost:7343`.

```
kg web          # start web server (default port 7343)
kg web --port 8080 # custom port
```

bash

Routes

Route	Content
/	Node index: all curated nodes + collapsible docs section
/node/<slug>	Node page: bullets, backlinks, related nodes (lazy-loaded)
/node/_doc-xxx	Source file page: metadata, chunks, syntax highlighting, GitHub link
/search?q=<query>	Search results page (FTS + vector blend)
/api/related/<slug>	JSON endpoint for lazy-loading related nodes

Table 9: Web viewer routes

Features

- **Backlinks:** every node page shows all nodes that reference `[[this-slug]]` in their bullets
- **Related nodes:** lazy-loaded sidebar using the node title + first 6 bullets as query
- **Source file rendering:** syntax highlighting (highlight.js CDN) and Markdown rendering (marked.js CDN)
- **Auto-linkification:** file paths like `src/kg/web.py` in bullet text are auto-linked to their `_doc-*` nodes
- **GitHub link:** doc pages show a “GitHub ↗” link pointing to the file at its last-modified commit hash

Troubleshooting

SQLite Disk I/O Error

On virtualized filesystems (OrbStack, NFS, Docker volumes), `graph.db` can become a 0-byte empty file, causing disk I/O error on all `kg` commands:

```
# Fix:  
kg stop  
rm .kg/index/graph.db*  
kg reindex      # stops watcher, rebuilds, restarts
```

bash

⚠ Warning

Always stop the watcher **before** removing `graph.db`. The watcher holds WAL/SHM files; deleting them while the watcher runs causes corruption on the next write.

Zero Files Indexed from Sources

If `kg status` shows 0 files indexed for a source:

- Check include patterns: use `**/*.py` not `*.py` (see Python 3.12 fnmatch note)
- If `use_git = true`: verify with `git ls-files` in the source directory that files are tracked
- Run `kg index` to trigger an immediate re-index cycle
- Check `kg status` config section for path/git validation errors

Embeddings Coverage Below 100%

After adding many nodes while the vector server was stopped, some nodes may have no embeddings. Start the server and recalibrate:

```
kg start      # ensure vector server is running  
kg calibrate # regenerates missing embeddings + recalibrates scores
```

bash

Vec Scores Missing from Calibration

If `kg status` shows “no vec calibration”, the vector server was likely stopped during the last `kg calibrate` run:

```
kg start      # start the vector server  
kg calibrate # re-run calibration
```

bash

Glossary

token budget: A credit counter per node tracking how many characters the node has contributed to context output. Triggers a review flag when credits-per-bullet exceeds the threshold. [12](#)

bullet: A single atomic fact, decision, gotcha, or task belonging to a node. Stored as a JSONL line with a stable UUID. [1](#)

calibration: The process of computing score quantile breakpoints from a random sample of nodes, used to normalize raw FTS and vector scores before fusion. [9](#)

MCP - Model Context Protocol: Open protocol for exposing tools to LLMs. kg exposes memory_context, memory_search, memory_show, and memory_add_bullet. [13](#)

node: A named collection of bullets identified by a slug. Stored as a directory under ` .kg/nodes/<slug>/` . [1](#)

reranker: A cross-encoder model (ms-marco-MiniLM-L-6-v2) that rescores the top-N hybrid search results using the full query-bullet pair. [9](#)

slug: A URL-safe lowercase identifier for a node, e.g. `my-topic` . Pattern: ` [a-z0-9][a-z0-9-]*` . [1](#)

watcher: A background process that uses inotify (Linux) to detect changes to ` .kg/nodes/` and immediately re-indexes modified nodes into SQLite. [2](#)